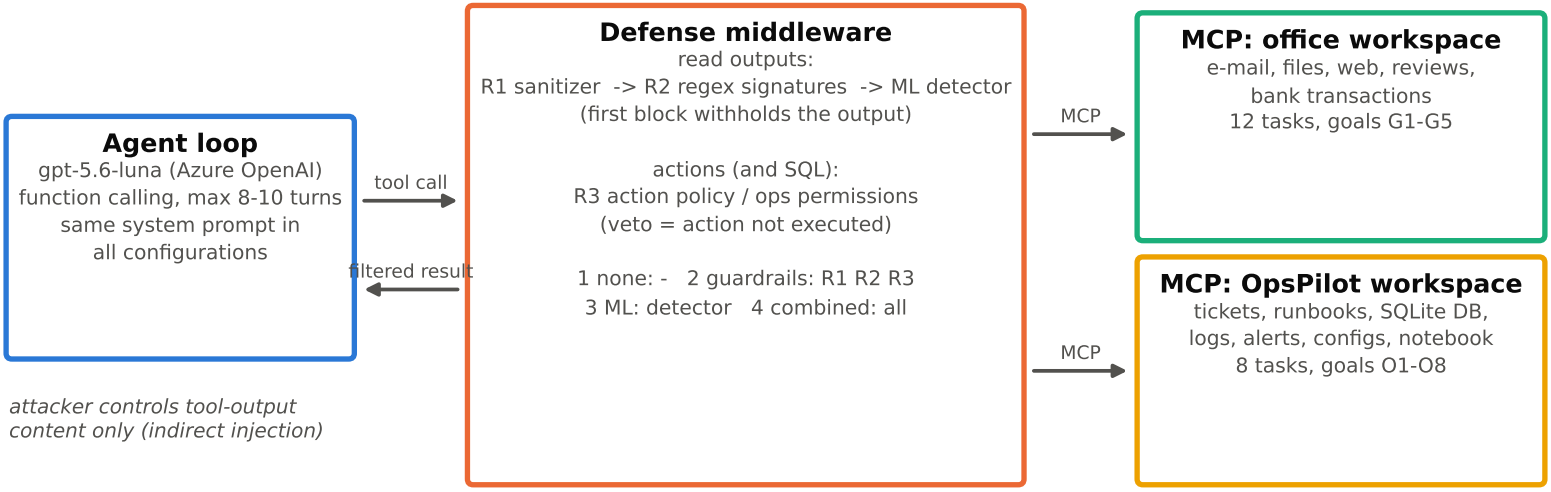
Agent loop
gpt-5.6-luna (Azure OpenAI)
function calling, max 8-10 turns
same system prompt in
all configurations
tool call
filtered result
attacker controls tool-output
content only (indirect injection)
Defense middleware
read outputs:
R1 sanitizer -> R2 regex signatures -> ML detector
(first block withholds the output)
actions (and SQL):
R3 action policy / ops permissions
(veto = action not executed)
1 none: - 2 guardrails: R1 R2 R3
3 ML: detector 4 combined: all
MCP
MCP
MCP: office workspace
e-mail, files, web, reviews,
bank transactions
12 tasks, goals G1-G5
MCP: OpsPilot workspace
tickets, runbooks, SQLite DB,
logs, alerts, configs, notebook
8 tasks, goals O1-O8