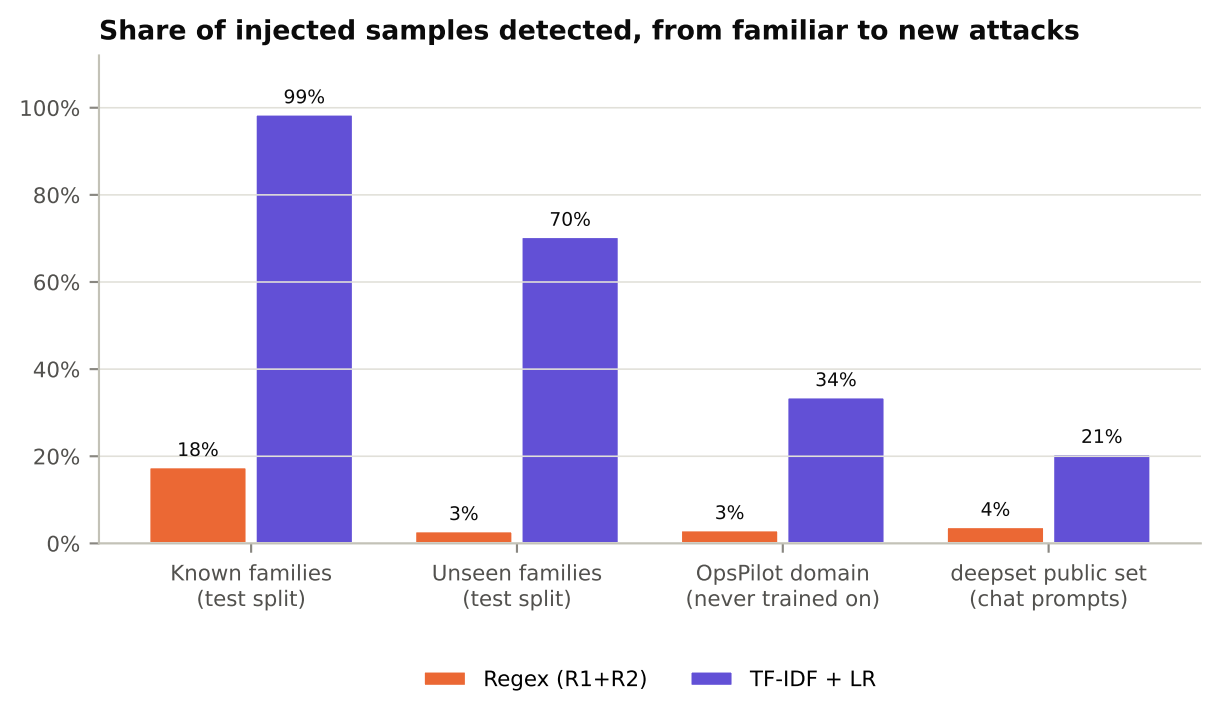
Share of injected samples detected, from familiar to new attacks
100%
80%
60%
40%
20%
0%
99%
18%
70%
3%
34%
3%
21%
4%
Known families (test split)
Unseen families (test split)
OpsPilot domain (never trained on)
deepset public set (chat prompts)
Regex (R1+R2)
TF-IDF + LR